

人體活動辨識模組

Human activity recognition module

國立台灣科技大學電子系
Department of Electronic and
Computer Engineering, NTUST






單元一：Cat & Dog分類CNN模型 實驗

Unit 1: Cat & Dog classification and CNN model experiment

國立台灣科技大學電子系

Department of Electronic and
Computer Engineering, NTUST





Unit 1: Cat & Dog classification and CNN model experiment

- What is machine learning?
- Perceptron
- Fully connected layer
- Convolution
- Transfer learning
- Lab 1





What is machine learning?


“Machine Learning”



f



“機器學習”



f



“Hello world”




f





“Cat”






What is machine learning?

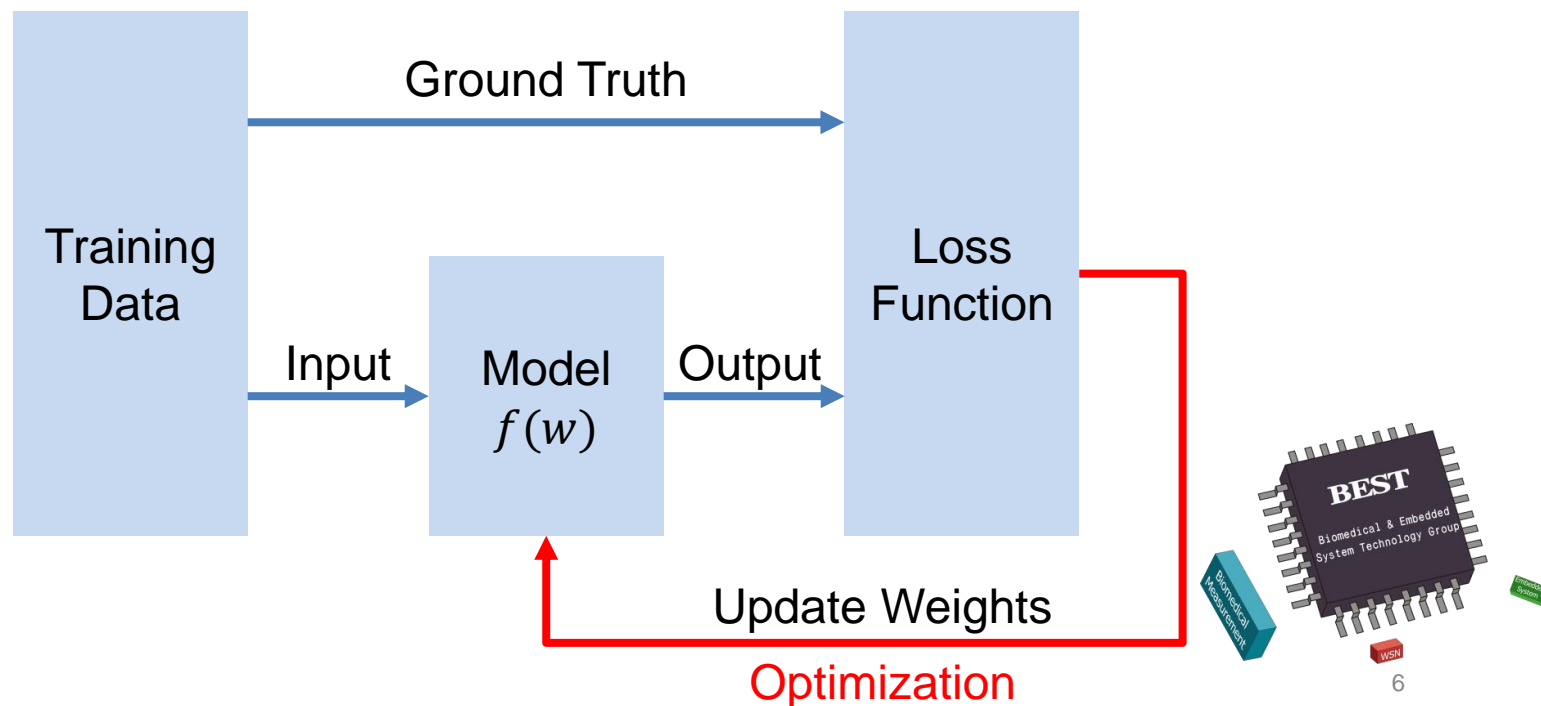
- Define a set of function $f(w_1, w_2, \dots, w_n)$.
- Get a best function $f(0, 1, \dots, 0)$.





Training



- Loss Function: Measure the difference between ground truth and model output.
- Optimization: Try to get better weights for the model.






Overfitting


- Overfitting means that prediction is too close to the training data.
- It may get a bad accuracy at unseen data.





Perceptron

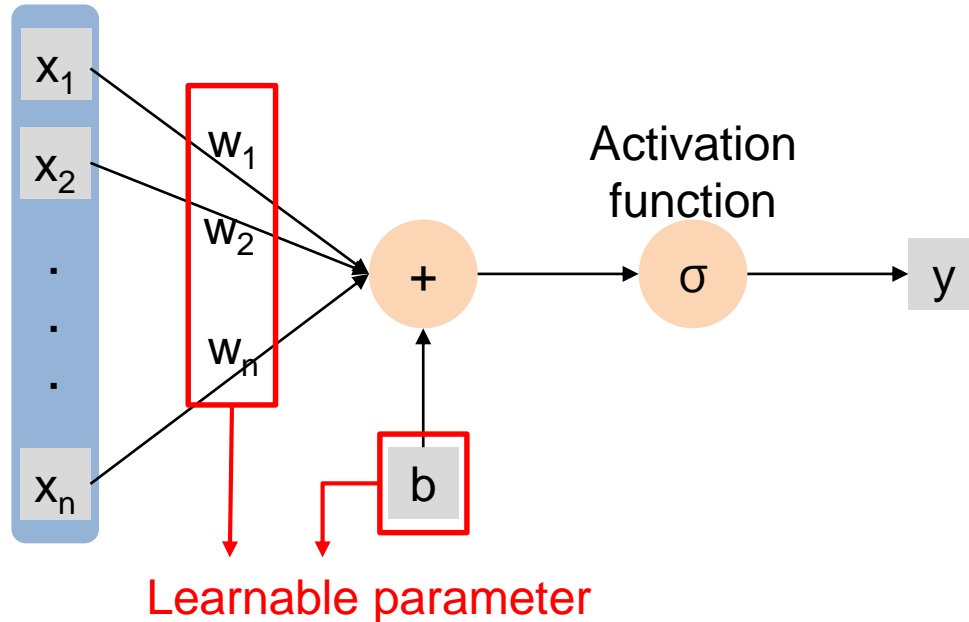
- Perceptron is a mathematical function modeled on the working of biological neurons.



Perceptron

- The output of perceptron is weighted sum of input and passed through a activation function.

Input



$$y = \sigma \left(\sum_{i=1}^n w_i x_i + b \right)$$





Activation function


- Activation function is a nonlinear function that mapping the result value in a range.

ex.

Sigmoid: $[-\infty, +\infty] \rightarrow [0, 1]$


Tanh : $[-\infty, +\infty] \rightarrow [-1, 1]$






Activation function

- Common activation function




$$\text{sigmoid} : \sigma = \frac{1}{1+e^{-x}}$$

limit output value
between 0 and 1




$$\text{tanh} : \sigma = \tanh(x)$$

limit output value
between -1 and 1



$$\text{ReLU} : \sigma = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

ignore negative
output






Perceptron example

- AND gate

$$w = [0.4, 0.5], b = -0.8, \sigma(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$



$$y = \sigma(0.4 * x_1 + 0.5 * x_2 - 0.8)$$


$$x = [0, 0], y = \sigma(0.4 * 0 + 0.5 * 0 - 0.8) = \sigma(-0.8) = 0$$

$$x = [0, 1], y = \sigma(0.4 * 0 + 0.5 * 1 - 0.8) = \sigma(-0.3) = 0$$

$$x = [1, 0], y = \sigma(0.4 * 1 + 0.5 * 0 - 0.8) = \sigma(-0.4) = 0$$

$$x = [1, 1], y = \sigma(0.4 * 1 + 0.5 * 1 - 0.8) = \sigma(+0.1) = 1$$






Perceptron example


- AND gate

Single level perceptron can learn only linearly separable patterns

| X_1 | X_2 | Q (label) |
|-------|-------|-----------|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 1 | 0 | 0 |
| 1 | 1 | 1 |

$$w = [0.4, 0.5], b = -0.8, \sigma(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$






Perceptron example

- XOR gate
 - Cannot find a suitable parameter for simulating XOR gate.

| X_1 | X_2 | Q (label) |
|-------|-------|-----------|
| 0 | 0 | 0 |
| 0 | 1 | 1 |
| 1 | 0 | 1 |
| 1 | 1 | 0 |






Perceptron example


- Single layer perceptron can learn only linear separable pattern.
- For more complicate problem, we need more perceptrons.






Perceptron example


- XOR gate
 - Idea: Combine more linear separable gates to perform XOR gate.






Perceptron example

- XOR gate



Fully connected layer

- Fully connected layer is a neuron network that all neurons are connected all the input neuron.




$$y_1 = \sigma(w_{11}x_1 + w_{21}x_2 + b_1)$$

$$y_2 = \sigma(w_{12}x_1 + w_{22}x_2 + b_2)$$

$$y_3 = \sigma(w_{13}x_1 + w_{23}x_2 + b_3)$$


$$y_4 = \sigma(w_{14}x_1 + w_{24}x_2 + b_4)$$





Fully connected layer

- If we cannot classify the training data, we may need more layers.





Convolution

- If we want to find a pattern: Vertical line in a image, we can use the below kernel (filter) to calculate the convolution.

| | | | | |
|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 1 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |

input

*

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 0 | 1 | 0 |
| 0 | 1 | 0 |

Kernel (filter)


=

| | | |
|---|--|--|
| 1 | | |
| | | |
| | | |

output

$$\begin{aligned} &1*0 + 0*1 + 0*0 + \\ &0*0 + 1*1 + 0*0 + \\ &0*0 + 0*1 + 1*0 \end{aligned}$$





Convolution

- Then we slide window and calculate the next value.

| | | | | |
|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 1 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |

input

*

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 0 | 1 | 0 |
| 0 | 1 | 0 |

Kernel (filter)


=

| | | |
|---|---|--|
| 1 | 1 | |
| | | |
| | | |

output

$$\begin{aligned} &0*0 + 0*1 + 0*0 + \\ &1*0 + 0*1 + 1*0 + \\ &0*0 + 1*1 + 0*0 \end{aligned}$$





Convolution

- Then we slide window and calculate the next value.

| | | | | |
|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 1 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |

input

*

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 0 | 1 | 0 |
| 0 | 1 | 0 |

Kernel (filter)

=

| | | |
|---|---|---|
| 1 | 1 | 1 |
| | | |
| | | |

output

$$\begin{aligned} &0*0 + 0*1 + 1*0 + \\ &0*0 + 1*1 + 0*0 + \\ &1*0 + 0*1 + 0*0 \end{aligned}$$





Convolution

- Then we slide window and calculate the next value.

| | | | | |
|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 1 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |

input

*

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 0 | 1 | 0 |
| 0 | 1 | 0 |

Kernel (filter)


=

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | | |
| | | |

output

$$\begin{aligned} &0*0 + 1*1 + 0*0 + \\ &0*0 + 0*1 + 1*0 + \\ &0*0 + 0*1 + 1*0 \end{aligned}$$





Convolution

| | | | | |
|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 1 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |

input

*

Kernel (filter)

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 0 | 1 | 0 |
| 0 | 1 | 0 |


=

$$\begin{aligned}
 &1*0 + 0*1 + 0*0 + \\
 &1*0 + 0*1 + 0*0 + \\
 &1*0 + 0*1 + 0*0
 \end{aligned}$$

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 2 | 1 |
| 0 | 3 | 0 |

output





Convolution

- The output indicates how many (vertical line) features the input has.

vertical line feature

| | | | | |
|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 1 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |

input

*

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 0 | 1 | 0 |
| 0 | 1 | 0 |

Kernel (filter)

=

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 2 | 1 |
| 0 | 3 | 0 |

output





Convolution


- Convolution operation can extract features on the image.
- But handcraft kernel is not efficient.
- So we use training data to train the kernel parameters.

| | | |
|----------|----------|----------|
| w_{11} | w_{12} | w_{13} |
| w_{21} | w_{22} | w_{23} |
| w_{31} | w_{32} | w_{33} |


Kernel (filter)

Learnable parameter





Convolution



| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | | | | | | | 0 |
| 0 | | | | | | | 0 |
| 0 | | | | | | | 0 |
| 0 | | | | | | | 0 |
| 0 | | | | | | | 0 |
| 0 | | | | | | | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

image with padding zero


| | | | | | | | |
|--|--|--|--|--|--|--|--|
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

image without padding zero





Convolution





stride = 2

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | | | | | | | 0 |
| 0 | | | | | | | 0 |
| 0 | | | | | | | 0 |
| 0 | | | | | | | 0 |
| 0 | | | | | | | 0 |
| 0 | | | | | | | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |




Convolution





Pooling

- Downsample the image to reduce the calculated complexity.






Transfer learning

- A pre-trained model is a saved network that was previously trained on a large dataset.
- You can customize the pre-trained model for your task.






Transfer learning


- VGG-16
 - The model accuracy is 92.7% on top-5 test of ImageNet, which is a dataset of over 14 million images belonging to 1000 classes.

| |
|-------------|
| Image |
| Conv 64 |
| Conv 64 |
| Max pooling |
| Conv 128 |
| Conv 128 |
| Max pooling |
| Conv 256 |
| Conv 256 |
| Max pooling |
| Conv 512 |
| Conv 512 |
| Conv 512 |
| Max pooling |
| Conv 512 |
| Conv 512 |
| Conv 512 |
| Max pooling |
| FC 4096 |
| FC 4096 |
| FC 1000 |



Transfer learning

- Task :
 - classify dog and cat
- Idea :
 - change last FC layer



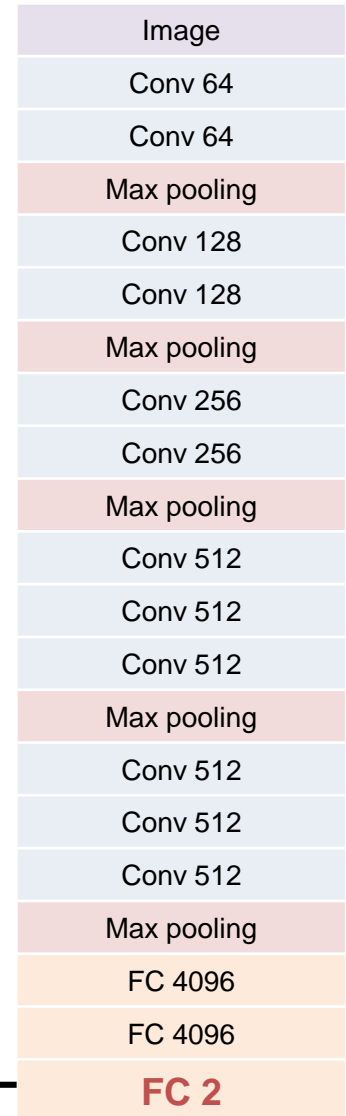



Transfer learning

- Try to use pre-trained weights when training in a small dataset.

pre-train from ImageNet

reinitialize and train





Transfer learning

- Use fewer pre-trained weights when training in a larger dataset.

pre-train from ImageNet

reinitialize and train


| |
|-------------|
| Image |
| Conv 64 |
| Conv 64 |
| Max pooling |
| Conv 128 |
| Conv 128 |
| Max pooling |
| Conv 256 |
| Conv 256 |
| Max pooling |
| Conv 512 |
| Conv 512 |
| Conv 512 |
| Max pooling |
| Conv 512 |
| Conv 512 |
| Conv 512 |
| Max pooling |
| FC 4096 |
| FC 4096 |
| FC 2 |



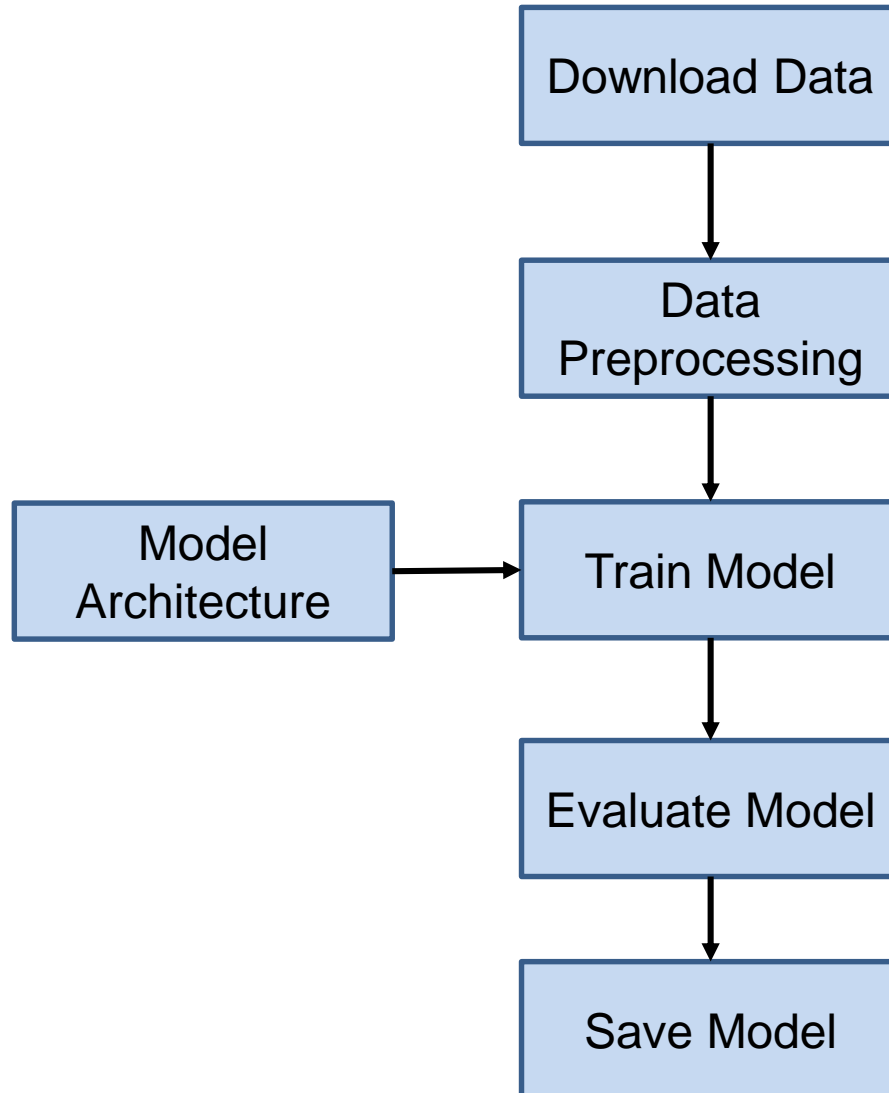
Lab 1


Cat & Dog classification and CNN model experiment





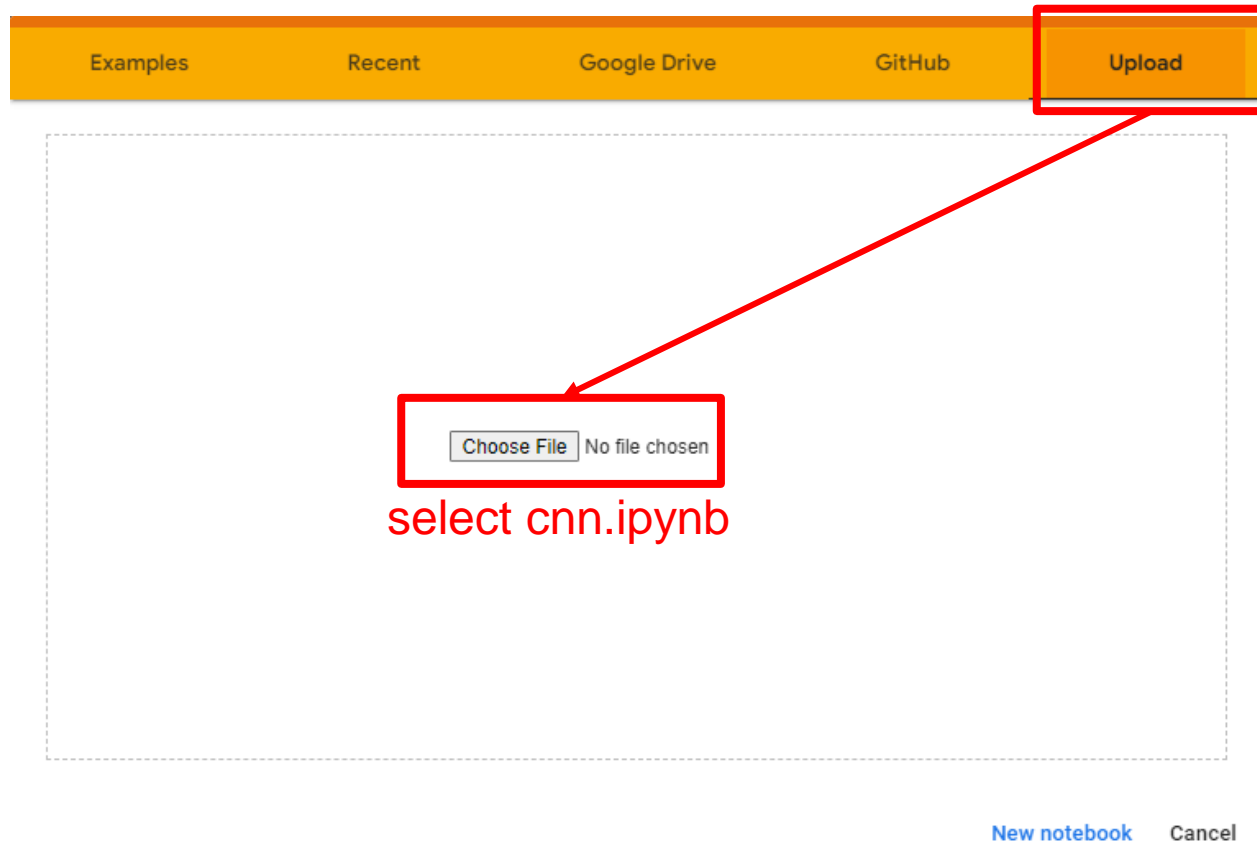
Flowchart





Open Colab

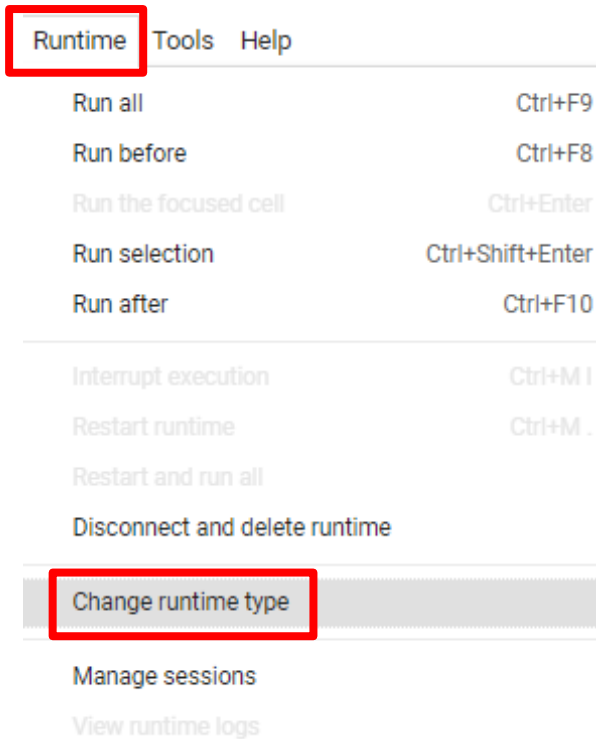
https://colab.research.google.com/?utm_source=scs-index





Colab setting

- Computing on a GPU in Colab.



Notebook settings


Hardware accelerator
GPU


To get the most out of Colab, avoid using a GPU unless you need one. [Learn more](#)

☐ Background execution

Want your notebook to keep running even after you close your browser? [Upgrade to Colab Pro+](#)

☐ Omit code cell output when saving this notebook







Download dataset

- The dataset is excerpted from the “Dogs vs. Cats” dataset on Kaggle.
- Train data : 2000 JPG of cats and dogs
- Validation data : 1000 JPG of cats and dogs

```
!wget --no-check-certificate \
  https://storage.googleapis.com/mledu-datasets/cats_and_dogs_filtered.zip \
  -O /tmp/cats_and_dogs_filtered.zip
```





Download dataset

- Unzip the file.

```
import zipfile


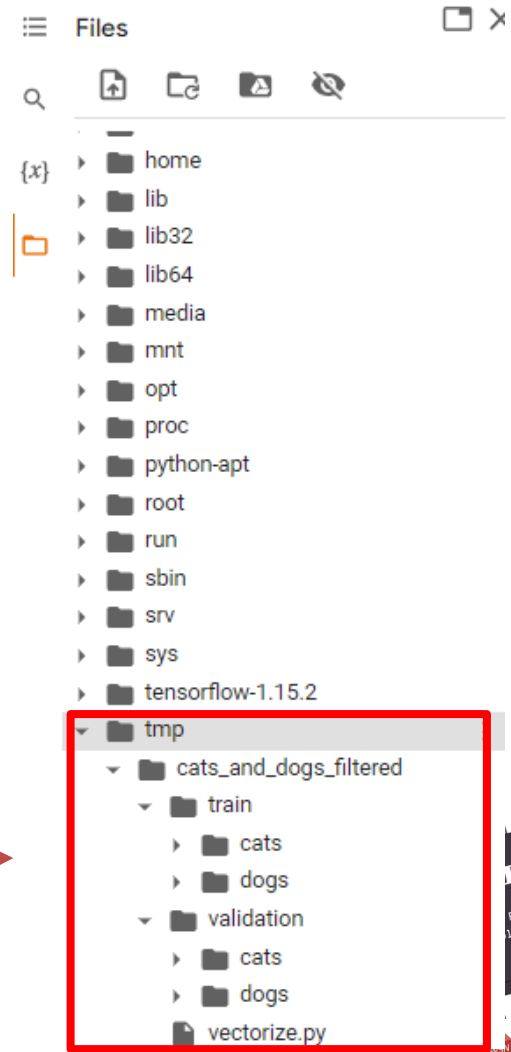

local_zip = '/tmp/cats_and_dogs_filtered.zip'
zip_ref = zipfile.ZipFile(local_zip, 'r')
zip_ref.extractall('/tmp')
zip_ref.close()
```






Download dataset

- Check the file hierarchy.
 - Go to the parent directory and check `tmp/cats_and_dogs_filtered`.





Data preprocessing

- Use `image_dataset_from_directory` to load image.

```
image_size = (224,224)
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    "/tmp/cats_and_dogs_filtered/train",
    seed=1337,
    image_size=image_size,
    label_mode="binary",
    batch_size=20
)
val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    "/tmp/cats_and_dogs_filtered/validation",
    seed=1337,
    image_size=image_size,
    label_mode="binary",
    batch_size=20
)
```






Data preprocessing

- Seed : random seed for shuffling data.
- Label_mode : 'binary' means that the labels are encoded as scalars with values 0 or 1.
- Image_size : resize the image.

```
image_size = (224,224)
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    "/tmp/cats_and_dogs_filtered/train",
    seed=1337,
    image_size=image_size,
    label_mode="binary",
    batch_size=20
)
```





Data preprocessing


- Rescale an input in the [0, 255] range to be in the [0, 1] range.

```
normalization_layer = tf.keras.layers.Rescaling(1./255)
train_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
val_ds = val_ds.map(lambda x, y: (normalization_layer(x), y))
```

- Check the value.

```
image_batch, labels_batch = next(iter(train_ds))
first_image = image_batch[0]
print(np.min(first_image), np.max(first_image))
```






Show image with label

```
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow((images[i].numpy()*255).astype("uint8"))
        plt.title(int(labels[i]))
        plt.axis("off")
```

0 : cat
1 : dog





Model architecture


- Import the layers you need.


```
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense  
from tensorflow.keras.layers import Activation, Flatten, Input
```

- Define input layer.

```
input_tensor=Input(shape=(224, 224, 3))
```

The shape of image





Model architecture

- Forward propagation the value by cascade layer.


```
input_tensor=Input(shape=(224, 224, 3))

x = Conv2D(64, (3, 3), activation='relu', padding='same')(input_tensor)
x = Conv2D(32, (3, 3), activation='relu', padding='same')(x)
x = MaxPooling2D((2, 2), strides=(2, 2))(x)

x = Conv2D(16, (3, 3), activation='relu', padding='same')(x)
x = Conv2D(16, (3, 3), activation='relu', padding='same')(x)
x = MaxPooling2D((2, 2), strides=(2, 2))(x)

x = Flatten()(x)
x = Dense(128, activation='relu')(x)
output_tensor = Dense(1, activation='sigmoid')(x)
model = Model(input_tensor, output_tensor)
```





Model architecture

```
x = Conv2D(16, (3, 3), activation='relu', padding='same')(x)
```

Number of kernel

Kernel size

One of “valid” or “same”.
“valid” means no padding.
“same” padding with zero.

Commonly used ‘relu’, ‘sigmoid’, ‘softmax’, ‘tanh’






Model architecture

```
x = MaxPooling2D((2, 2), strides=(2, 2))(x)
```

Pool size

Specifies how far the pooling window moves for each pooling step.






Model architecture

- Change the 3d array to 1d array.


```
x = Flatten() (x)
```



- Dense means fully connected layer.

```
x = Dense(128, activation='relu') (x)
```

Number of neuron






Check your model

```
model.summary()
```

| Layer (type) | Output Shape | Param # |
|---------------------------------|-----------------------|---------|
| input_11 (InputLayer) | [(None, 224, 224, 3)] | 0 |
| conv2d_25 (Conv2D) | (None, 224, 224, 64) | 1792 |
| conv2d_26 (Conv2D) | (None, 224, 224, 32) | 18464 |
| max_pooling2d_11 (MaxPooling2D) | (None, 112, 112, 32) | 0 |
| conv2d_27 (Conv2D) | (None, 112, 112, 16) | 4624 |
| conv2d_28 (Conv2D) | (None, 112, 112, 16) | 2320 |
| max_pooling2d_12 (MaxPooling2D) | (None, 56, 56, 16) | 0 |
| flatten_9 (Flatten) | (None, 50176) | 0 |
| dense_18 (Dense) | (None, 128) | 6422656 |
| dense_19 (Dense) | (None, 1) | 129 |
| Total params: 6,449,985 | | |
| Trainable params: 6,449,985 | | |
| Non-trainable params: 0 | | |





Train model


- Configure the model, set optimizer and loss function.

```
adam = tf.keras.optimizers.Adam(learning_rate=0.001)
momentum = tf.keras.optimizers.SGD(learning_rate=0.001, momentum=0.9)

model.compile(optimizer=momentum, loss='binary_crossentropy', metrics=['accuracy'])
```

Now is “momentum”





Train model

- Train the model and store record of training loss values and metrics values.
- Epochs : Number of epochs to train the model.


```
hist = model.fit(train_ds,  
                 epochs=100,  
                 validation_data=val_ds,  
                 verbose=1)
```






Train model

```
Epoch 1/100
100/100 [=====] - 111s 709ms/step - loss: 0.6932 - accuracy: 0.4975 - val_loss: 0.6931 - val_accuracy: 0.5000
Epoch 2/100
100/100 [=====] - 71s 707ms/step - loss: 0.6931 - accuracy: 0.4935 - val_loss: 0.6930 - val_accuracy: 0.5830
Epoch 3/100
100/100 [=====] - 71s 708ms/step - loss: 0.6930 - accuracy: 0.5470 - val_loss: 0.6929 - val_accuracy: 0.5110
Epoch 4/100
100/100 [=====] - 71s 707ms/step - loss: 0.6931 - accuracy: 0.4970 - val_loss: 0.6929 - val_accuracy: 0.5040
Epoch 5/100
100/100 [=====] - 71s 708ms/step - loss: 0.6929 - accuracy: 0.5245 - val_loss: 0.6928 - val_accuracy: 0.5050
Epoch 6/100
100/100 [=====] - 71s 707ms/step - loss: 0.6928 - accuracy: 0.5485 - val_loss: 0.6927 - val_accuracy: 0.5210
Epoch 7/100
100/100 [=====] - 71s 708ms/step - loss: 0.6929 - accuracy: 0.4985 - val_loss: 0.6926 - val_accuracy: 0.5160
Epoch 8/100
100/100 [=====] - 71s 708ms/step - loss: 0.6928 - accuracy: 0.5240 - val_loss: 0.6926 - val_accuracy: 0.5040
Epoch 9/100
100/100 [=====] - 71s 708ms/step - loss: 0.6926 - accuracy: 0.5520 - val_loss: 0.6924 - val_accuracy: 0.5240
Epoch 10/100
100/100 [=====] - 71s 707ms/step - loss: 0.6926 - accuracy: 0.5330 - val_loss: 0.6923 - val_accuracy: 0.6050
Epoch 11/100
100/100 [=====] - 71s 708ms/step - loss: 0.6926 - accuracy: 0.5430 - val_loss: 0.6922 - val_accuracy: 0.5710
Epoch 12/100
100/100 [=====] - 71s 707ms/step - loss: 0.6923 - accuracy: 0.5530 - val_loss: 0.6920 - val_accuracy: 0.5580
Epoch 13/100
100/100 [=====] - 71s 708ms/step - loss: 0.6921 - accuracy: 0.5530 - val_loss: 0.6918 - val_accuracy: 0.5400
Epoch 14/100
100/100 [=====] - 71s 707ms/step - loss: 0.6921 - accuracy: 0.5455 - val_loss: 0.6917 - val_accuracy: 0.5220
Epoch 15/100
100/100 [=====] - 71s 708ms/step - loss: 0.6920 - accuracy: 0.5250 - val_loss: 0.6913 - val_accuracy: 0.5720
Epoch 16/100
100/100 [=====] - 71s 708ms/step - loss: 0.6916 - accuracy: 0.5495 - val_loss: 0.6910 - val_accuracy: 0.5720
Epoch 17/100
100/100 [=====] - 71s 708ms/step - loss: 0.6911 - accuracy: 0.5645 - val_loss: 0.6907 - val_accuracy: 0.5480
Epoch 18/100
100/100 [=====] - 71s 708ms/step - loss: 0.6911 - accuracy: 0.5465 - val_loss: 0.6905 - val_accuracy: 0.5330
```





Evaluate model

```
import matplotlib.pyplot as plt
from keras.preprocessing.image import img_to_array, load_img

# Retrieve a list of accuracy results on training and test data
# sets for each training epoch
acc = hist.history['accuracy']
val_acc = hist.history['val_accuracy']


# Retrieve a list of list results on training and test data
# sets for each training epoch
loss = hist.history['loss']
val_loss = hist.history['val_loss']


# Get number of epochs
epochs = range(len(acc))

# Plot training and validation accuracy per epoch
plt.plot(epochs, acc)
plt.plot(epochs, val_acc)
plt.title('Training and validation accuracy')

plt.figure()




# Plot training and validation loss per epoch
plt.plot(epochs, loss)
plt.plot(epochs, val_loss)
plt.title('Training and validation loss')
```





Evaluate model

- As you can see, the model is overfitting.
- Our training accuracy gets close to 100%, but validation accuracy stalls at 60%.






Overfitting

- Overfitting is a tough problem when we train the model.
- Overfitting happens when a model learns the detail and noise in the training data to the extent that it negatively impacts the performance of the model on new data.






Save / load model

- Save your model to deploy on different platform.
- You can find your file in /content/myModel.h5.

```
model.save("myModel.h5")  
  
new_model = tf.keras.models.load_model('myModel.h5')  
  
new_model.evaluate(val_ds, verbose=1)
```

```
50/50 [=====] - 11s 214ms/step - loss: 0.0730 - accuracy: 0.9770  
[0.07302631437778473, 0.9769999980926514]
```



Inference

```
animal = "cat" # cat or dog
index = 2222 # 2000-2499
img_path = f"/tmp/cats_and_dogs_filtered/validation/{animal}s/{animal}.{index}.jpg"
img = load_img(img_path, target_size=(224, 224)) # this is a PIL image
plt.title(img_path)
plt.axis("off")
plt.imshow(img)

x = img_to_array(img) # Numpy array with shape (224, 224, 3)
x = x.reshape((1,) + x.shape) # Numpy array with shape (1, 224, 224, 3)
x /= 255 # Rescale by 1/255


import time
start = time.time() # record start time

result = new_model.predict(x) # inference(predict)
if result < 0.5:
    print("It is a cat.")
else:
    print("It is a dog.")
finish = time.time() # record finish time

print("Result = %f" %result)
print("Test time :%f second." %(finish-start))
```

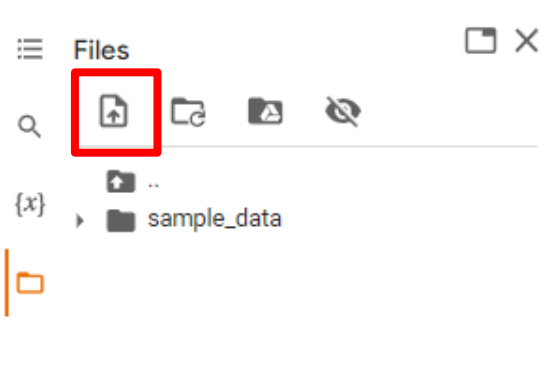
```
It is a cat.
Result = 0.004889
Test time :0.167567 second.
/tmp/cats_and_dogs_filtered/validation/cats/cat.2222.jpg
```





Upload/load model

- Upload pre-trained model and evaluate the model.




```
model = tf.keras.models.load_model("myModel.h5")
```

```
model.evaluate(val_ds, verbose=1)
```

```
50/50 [=====] - 494s 10s/step - loss: 0.0730 - accuracy: 0.9770  
[0.07302630692720413, 0.9769999980926514]
```





Assignment 1

- Select two of the models in the right chart.
- Implement and compare their accuracies.

| A | B | C | D | E |
|-----------------------------------|----------|----------|----------|----------|
| conv3-4 | conv3-4 | conv3-4 | conv3-4 | conv3-4 |
| | conv3-4 | conv3-4 | conv3-4 | conv3-4 |
| maxpool(size=(2,2),strides=(2,2)) | | | | |
| conv3-8 | conv3-8 | conv3-8 | conv3-8 | conv3-8 |
| | conv3-8 | conv3-8 | conv3-8 | conv3-8 |
| maxpool(size=(2,2),strides=(2,2)) | | | | |
| conv3-16 | conv3-16 | conv3-16 | conv3-16 | conv3-16 |
| conv3-16 | conv3-16 | conv3-16 | conv3-16 | conv3-16 |
| | | conv3-16 | conv3-16 | conv3-16 |
| | | | conv3-16 | conv3-16 |
| maxpool(size=(2,2),strides=(2,2)) | | | | |
| conv3-32 | conv3-32 | conv3-32 | conv3-32 | conv3-32 |
| conv3-32 | conv3-32 | conv3-32 | conv3-32 | conv3-32 |
| | | conv3-32 | conv3-32 | conv3-32 |
| | | | conv3-32 | conv3-32 |
| maxpool(size=(2,2),strides=(2,2)) | | | | |
| conv3-64 | conv3-64 | conv3-64 | conv3-64 | conv3-64 |
| conv3-64 | conv3-64 | conv3-64 | conv3-64 | conv3-64 |
| | | conv3-64 | conv3-64 | conv3-64 |
| | | | conv3-64 | conv3-64 |
| maxpool(size=(2,2),strides=(2,2)) | | | | |
| FC-128 | | | | |
| FC-64 | | | | |
| FC-1 | | | | |
| sigmoid | | | | |

ConvA-B => kernel size: (A,A), number of the kernel: B

FC-N => number of neuron: N






Assignment 1

- Finish a report.
 - What is observed & learned?
 - Screen dump:
 - `Model.summary()`
 - Plot training and validation accuracy
 - Plot training and validation loss






Submission

- Hierarchy:
 - Model_1.ipynb
 - Model_2.ipynb
 - Report.pdf
- Compress all above files in a single zip file named StudentID_lab1.zip.





References

- K. Simonyan and A. Zisserman, “Very deep convolutional networks for large-scale image recognition,” arXiv:1409.1556, 2014.
- Stanford University CS231n, “Deep Learning for Computer Vision.” [Online]. Available: <http://cs231n.stanford.edu/>.
- H. Y. Lee, “Hung-yi Lee youtube channel” [Online]. Available: <https://www.youtube.com/c/HungyiLeeNTU>.
- Kaggle, “Dogs vs. Cats.” [Online]. Available: <https://www.kaggle.com/c/dogs-vs-cats>.

